

# Brave new world of unified cgroups

Michał Sekletár  
Zbigniew Jędrzejewski-Szmek



*msekleta@redhat.com*  
*zbyszek@in.waw.pl*



Brno, 2020.01.26

Why this talk?

# Why this talk?

Unified hierarchy (a.k.a.) cgroups v2 is the default default in systemd 243 (Sept. 2019)

<https://fedoraproject.org/wiki/Changes/CGroupsV2> (for F31+)

Version two is just nicer!

# Control groups

Control groups (cgroups) is a Linux subsystem that has two main purposes:

- ▶ Process tracking
- ▶ Resource distribution

# Control groups

**Cgroup** – associates a set of tasks with a set of parameters for one or more controllers

**Controller** – schedules a resource or applies per-cgroup limits


**Hierarchy** – set of cgroups arranged in a tree, every process is in exactly one cgroup

## Brief history of cgroups in the linux kernel

2008	kernel 2.6.24	task control groups
2012		v2 rawrite announced
2013		v2 available
2015	4.3	prols controller
2016	4.5	v2 stable!
2017	4.14 4.15	threaded mode CPU controller
2019	5.0 5.1	cpu set controller freezer

Why?

# What's wrong with v1





# What's wrong with v1

- ▶ design follows implementation

# What's wrong with v1

- ▶ design follows implementation
- ▶ inconsistent interface

# What's wrong with v1

- ▶ design follows implementation
- ▶ inconsistent interface
- ▶ infinite number of hierarchies

# What's wrong with v1

- ▶ design follows implementation
- ▶ inconsistent interface
- ▶ infinite number of hierarchies
- ▶ not hierarchical

# What's wrong with v1

- ▶ design follows implementation
- ▶ inconsistent interface
- ▶ infinite number of hierarchies
- ▶ not hierarchical
- ▶ unusable limits

# What's wrong with v1

- ▶ design follows implementation
- ▶ inconsistent interface
- ▶ infinite number of hierarchies
- ▶ not hierarchical
- ▶ unusable limits
- ▶ no cooperation between controllers

# What's wrong with v1

- ▶ design follows implementation
- ▶ inconsistent interface
- ▶ infinite number of hierarchies
- ▶ not hierarchical
- ▶ unusable limits
- ▶ no cooperation between controllers
- ▶ threads not processes

# What's wrong with v1

- ▶ design follows implementation
- ▶ inconsistent interface
- ▶ infinite number of hierarchies
- ▶ not hierarchical
- ▶ unusable limits
- ▶ no cooperation between controllers
- ▶ threads not processes
- ▶ no secure delegation



```
$ ls /sys/fs/cgroup/memory/
```

```
...
```

```
memory.limit_in_bytes          memory.kmem.tcp.limit_in_bytes
memory.usage_in_bytes          memory.kmem.tcp.usage_in_bytes
memory.max_usage_in_bytes      memory.kmem.tcp.max_usage_in_bytes
memory.soft_limit_in_bytes
```

```
memory.kmem.limit_in_bytes     memory.memsw.limit_in_bytes
memory.kmem.usage_in_bytes     memory.memsw.usage_in_bytes
memory.kmem.max_usage_in_bytes memory.memsw.max_usage_in_bytes
```

```
memory.kmem.slabinfo           memory.use_hierarchy
memory.move_charge_at_immigrate cgroup.sane_behavior
```

```
...
```

## inconsistent interface

v1	default	range
cpu.shares	1024	2-262144
blkio.bfq.weight	500	10-1000
v2		
cpu.weight	100	1-10000
io.weight <sup>1</sup>	100	1-10000

---

<sup>1</sup><https://github.com/systemd/systemd/pull/13335>

## Design:

- ▶ single hierarchy
- ▶ consistent interface
- ▶ small number of controllers: memory, io, pids, cpu, cpuset
- ▶ controllers are fully hierarchical
- ▶ (controllers can be turned off midway through the tree)
- ▶ high-level knobs
- ▶ soft limits
- ▶ non-hierarchical controllers are gone

## Old vs. New

v1 controller	v2 solution
memory	memory
cpu+cpuacct	cpu
cpuset	cpuset
blkio	io
pids	pids
hugetlb	hugetlb (kernel 5.6)

## Old vs. New

v1 controller	v2 solution
memory	memory
cpu+cpuacct	cpu
cpuset	cpuset
blkio	io
pids	pids
hugetlb	hugetlb (kernel 5.6)
freezer	replaced by cgroup.freeze

## Old vs. New

v1 controller	v2 solution
memory	memory
cpu+cpuacct	cpu
cpuset	cpuset
blkio	io
pids	pids
hugetlb	hugetlb (kernel 5.6)
freezer	replaced by cgroup.freeze
devices	eBPF filters

## Old vs. New

v1 controller	v2 solution
memory	memory
cpu+cpuacct	cpu
cpuset	cpuset
blkio	io
pids	pids
hugetlb	hugetlb (kernel 5.6)
freezer	replaced by cgroup.freeze
devices	eBPF filters
net_cls,net_prio	matching by cgroup, eBPF
perf_event	eBPF

# Delegation

- ▶ Less-privileged process owns a part of the cgroup tree
- ▶ Implemented through file system permissions
- ▶ Delegate=yes in systemd units
- ▶ Cutoff is not at the directory level



```
$ sudo systemd-cgls
```

```
Control group /:  
- .slice  
├─ user.slice  
│   └─ user-6.slice  
│       └─ user@6.service ...  
│           └─ init.scope  
│               ├── 1963 /usr/lib/systemd/systemd --user  
│               └── 2001 (sd-pam)  
└─ user-1000.slice  
    └─ user@1000.service ...  
        ├── gsd-xsettings.service  
        │   └── 412129 /usr/libexec/gsd-xsettings  
        ├── gvfs-goa-volume-monitor.service  
        │   └── 412027 /usr/libexec/gvfs-goa-volume-monitor  
        ├── gsd-power.service  
        │   └── 412104 /usr/libexec/gsd-power  
        ├── dbus\x2d:1.1\x2dorg.gnome.Epiphany.SearchProvider.slice  
        │   └── dbus-:1.1-org.gnome.Epiphany.SearchProvider@0.service  
        │       └── 415659 /usr/libexec/epiphany-search-provider  
        ├── xdg-permission-store.service  
        │   └── 411997 /usr/libexec/xdg-permission-store  
        ├── dbus-broker.service  
        │   ├── 411532 /usr/bin/dbus-broker-launch --scope user  
        │   └── 411533 dbus-broker --log 4 --controller 11 --machine-id 08a5690a2eed47cf92ac0a  
        ├── xdg-document-portal.service  
        │   └── 412300 /usr/libexec/xdg-document-portal  
        ├── dbus\x2d:1.1\x2dorg.gnome.OnlineAccounts.slice  
        │   └── dbus-:1.1-org.gnome.OnlineAccounts@0.service  
        │       └── 412024 /usr/libexec/goa-daemon  
        ├── tracker-store.service  
        │   └── 468488 /usr/libexec/tracker-store  
        └── .
```

```
$ ls -l .../user.slice/user-1000.slice/user@1000.service
-r--r--r--. root      root      cgroup.controllers
-r--r--r--. root      root      cgroup.events
-rw-r--r--. root      root      cgroup.freeze
-rw-r--r--. root      root      cgroup.max.depth
-rw-r--r--. root      root      cgroup.max.descendants
-r--r--r--. root      root      cgroup.stat
-rw-r--r--. zbyszek   zbyszek   cgroup.procs
-rw-r--r--. zbyszek   zbyszek   cgroup.threads
-rw-r--r--. zbyszek   zbyszek   cgroup.subtree_control
-r--r--r--. root      root      pids.current
-r--r--r--. root      root      pids.events
-rw-r--r--. root      root      pids.max
...
drwxr-xr-x. zbyszek   zbyszek   pipewire.service/
drwxr-xr-x. zbyszek   zbyszek   pulseaudio.service/
drwxr-xr-x. zbyszek   zbyszek   xdg-desktop-portal-gtk.service/
drwxr-xr-x. zbyszek   zbyszek   xdg-desktop-portal.service/
drwxr-xr-x. zbyszek   zbyszek   xdg-document-portal.service/
```

# Delegation

is not all or nothing

- ▶ Delegate=io pids memory ...
- ▶ Delegation may be nested
- ▶ Resources are divided hierarchically

# Threaded mode

## Threaded mode

v1 operated on threads, v2 operates on processes

# Threaded mode

v1 operated on threads, v2 operates on processes

Only supported by selected controllers (cpu, cpuset)

# Threaded mode

v1 operated on threads, v2 operates on processes

Only supported by selected controllers (cpu, cpuset)

A leaf cgroup may be switched to threaded mode and further subdivided

```
$ echo threaded > $CGROUP/cgroup.type
```

# Threaded mode

v1 operated on threads, v2 operates on processes

Only supported by selected controllers (cpu, cpuset)

A leaf cgroup may be switched to threaded mode and further subdivided

```
$ echo threaded > $CGROUP/cgroup.type
```

cgroup.threads



# Threaded mode

v1 operated on threads, v2 operates on processes

Only supported by selected controllers (cpu, cpuset)

A leaf cgroup may be switched to threaded mode and further subdivided

```
$ echo threaded > $CGROUP/cgroup.type
```

cgroup.threads

Used by libvirt to manage QEMU vCPU threads

# Resource management – Resource distribution models

## ▶ **Weights**

- ▶ Resource is distributed by adding up the weights of all sub-cgroups and giving each the fraction matching its ratio against the sum.
- ▶ Usually used to distribute stateless resources (CPU time)
- ▶ Example: `cpu.weight` ([1-10000], default 100)

# Resource management – Resource distribution models

## ▶ **Weights**

- ▶ Resource is distributed by adding up the weights of all sub-cgroups and giving each the fraction matching its ratio against the sum.
- ▶ Usually used to distribute stateless resources (CPU time)
- ▶ Example: `cpu.weight` ([1-10000], default 100)

## ▶ **Limits**

- ▶ Cgroup can consume up to configured amount of the resource
- ▶ Overcommit is allowed (i.e. sum of sub-cgroup limits can exceed limit of the parent cgroup)
- ▶ Example: `memory.max`

# Resource management – Resource distribution models

## ▶ **Weights**

- ▶ Resource is distributed by adding up the weights of all sub-cgroups and giving each the fraction matching its ratio against the sum.
- ▶ Usually used to distribute stateless resources (CPU time)
- ▶ Example: `cpu.weight` ([1-10000], default 100)

## ▶ **Limits**

- ▶ Cgroup can consume up to configured amount of the resource
- ▶ Overcommit is allowed (i.e. sum of sub-cgroup limits can exceed limit of the parent cgroup)
- ▶ Example: `memory.max`

## ▶ **Protections**

- ▶ Cgroup is protected (but not guaranteed) up to configured amount of the resource
- ▶ Overcommit is also allowed
- ▶ Example: `memory.low`

# Resource management – Resource distribution models

## ▶ **Weights**

- ▶ Resource is distributed by adding up the weights of all sub-cgroups and giving each the fraction matching its ratio against the sum.
- ▶ Usually used to distribute stateless resources (CPU time)
- ▶ Example: `cpu.weight` ([1-10000], default 100)

## ▶ **Limits**

- ▶ Cgroup can consume up to configured amount of the resource
- ▶ Overcommit is allowed (i.e. sum of sub-cgroup limits can exceed limit of the parent cgroup)
- ▶ Example: `memory.max`

## ▶ **Protections**

- ▶ Cgroup is protected (but not guaranteed) up to configured amount of the resource
- ▶ Overcommit is also allowed
- ▶ Example: `memory.low`

## ▶ **Allocations**

- ▶ Exclusive allocations of the absolute amount of a finite resource (e.g. real-time budget)

# Memory protections and limits

LIMITs HARD

---

oom killer

SOFT

---

throttling  
reclaim pressure

PROTECTIONS LOW

---

😊

MIN

---

reclaim only if no  
un-protected

no reclaim  
(oom killer instead)

## Resource management – Memory

cgroup v2 memory controller implements memory usage limits but also very useful protections (not available in cgroup v1),

- ▶ **memory.min** – Hard memory protection. If memory usage is below the limit the cg memory won't be reclaimed.
- ▶ **memory.low** – Soft memory protection. If memory usage is below the limit the cg memory can be reclaimed only if there is no memory to be reclaimed from unprotected cgroups.
- ▶ **memory.high** – Memory throttle limit. If memory usage goes above the limit the processes in the cgroup are throttled and put under heavy reclaim pressure.
- ▶ **memory.max** – Hard limit for memory usage (OOM killer is invoked if usage goes above the limit)

## Resource management – Block I/O

`io` controller in `cgroup v2` allows for very fine grained settings (per device)

- ▶ **`io.weight`** – IO weight (work-conserving)
- ▶ **`io.max`** – per-device absolute bandwidth limit (e.g. 8:16  
`rbps=2097152 wiops=120`)
- ▶ **`io.latency`** – per-device latency target protection (e.g. 8:16  
`target=10000`)




# Status quo

v1 only: k8s, CRI-O, Docker, Containerd, runc (in progress),  
OCI runtime spec

v2 too: Buildah+Podman+skopeo, crun, libvirt, JVM, snapd,  
systemd

# Status of container runtimes

## Which components support unified hierarchy?



# Summary

control groups v2:

- ▶ Fully hierarchical with safe delegation
- ▶ Consistent interface
- ▶ Efficient and scalable notifications
- ▶ Fewer controllers, high-level knobs
- ▶ Soft limits
- ▶ eBPF!
- ▶ Better monitoring: Pressure Stall Information (PSI)

# Links

this talk: <https://github.com/keszybz/cgroupsv2/raw/master/cgroupsv2.pdf>

docs:

<https://www.kernel.org/doc/html/latest/admin-guide/cgroup-v1/>

<https://www.kernel.org/doc/html/latest/admin-guide/cgroup-v2.html>

<https://facebookmicrosites.github.io/cgroup2/docs/overview>

[systemd.resource-control\(5\)](#)

[https://systemd.io/CGROUP\\_DELEGATION.html](https://systemd.io/CGROUP_DELEGATION.html)

recent changes:

<https://www.redhat.com/sysadmin/fedora-31-control-group-v2>

<https://fedoraproject.org/wiki/Changes/CGroupsV2>

<https://www.youtube.com/watch?v=GznkuTXq8AQ&t=1s>

<https://medium.com/nttlabs/cgroup-v2-596d035be4d7>

<https://www.youtube.com/watch?v=yZpNsDe4Qzg> (Michael Kerrisk)

freezer for cgroup v2 v5.1-rc3-45-g76f969e894

<https://lwn.net/Articles/772377/>

[https://bugzilla.redhat.com/show\\_bug.cgi?id=1727149](https://bugzilla.redhat.com/show_bug.cgi?id=1727149) libvirt support in 5.5.0

[https://bugzilla.redhat.com/show\\_bug.cgi?id=1438079](https://bugzilla.redhat.com/show_bug.cgi?id=1438079) snapd support in snapd-2.41-1.fc31

<https://github.com/opencontainers/runc/pull/2113> for libcontainer

<https://github.com/opencontainers/runc/issues/654> for runc

<https://github.com/kubernetes/enhancements/pull/1370/files> for k8s

[codesearch.debian.net/search?q=cgroup.type](https://codesearch.debian.net/search?q=cgroup.type)

<https://www.kernel.org/doc/html/latest/accounting/psi.html>

# Links

history:

[https://kernelnewbies.org/Linux\\_2\\_6\\_24#Task\\_Control\\_Groups](https://kernelnewbies.org/Linux_2_6_24#Task_Control_Groups)

[https://kernelnewbies.org/Linux\\_3.16#Unified\\_Control\\_Group\\_hierarchy](https://kernelnewbies.org/Linux_3.16#Unified_Control_Group_hierarchy)

State of CPU controller in cgroup v2 (2016)

LWN: A milestone for control groups (2017)

Linux 4.15

v4.14-rc2-7-g0d5936344f

<https://www.youtube.com/watch?v=PzpG40WiEfM>

[https://www.youtube.com/watch?v=ikZ8\\_mRotT4](https://www.youtube.com/watch?v=ikZ8_mRotT4)